


PaaS

Heroku

CCS3341 Cloud Computing


Dr S. Veloudis

PaaS

- Platform as a Service
- But what is a platform 
 - A platform here refers to a software framework through which software solutions can be developed
 - It provides a complete operational and development environment for a multitude of programming languages
 - It offers core software functionality for developing a cloud application which would otherwise need to be engineered from scratch
 - A cloud platform therefore serves as a launch pad for cloud software
 - PaaS offerings also include the infrastructure on which cloud applications are deployed
 - This infrastructure may be provided by a 3rd party IaaS provider

Note: A PaaS provider may at the same time be an IaaS consumer!

- PaaS vs IaaS



source: blogs.technet.microsoft.com

PaaS

- Cloud platforms are offered as a service
 - No need to deploy and maintain the platform on proprietary equipment
 - The PaaS user simply codes up a software solution and doesn't worry about platform deployment details
- PaaS abstracts away interaction with the 'bare metal'
 - Frees developers from the burden of administering the OS
- PaaS customers are charged for the infrastructure on which the PaaS-developed cloud application is deployed
 - Clearly, the more resources this infrastructure is equipped with, the higher the charges



- Supports a multitude of programming languages



Node.js



Ruby



Java



PHP



Python



Go



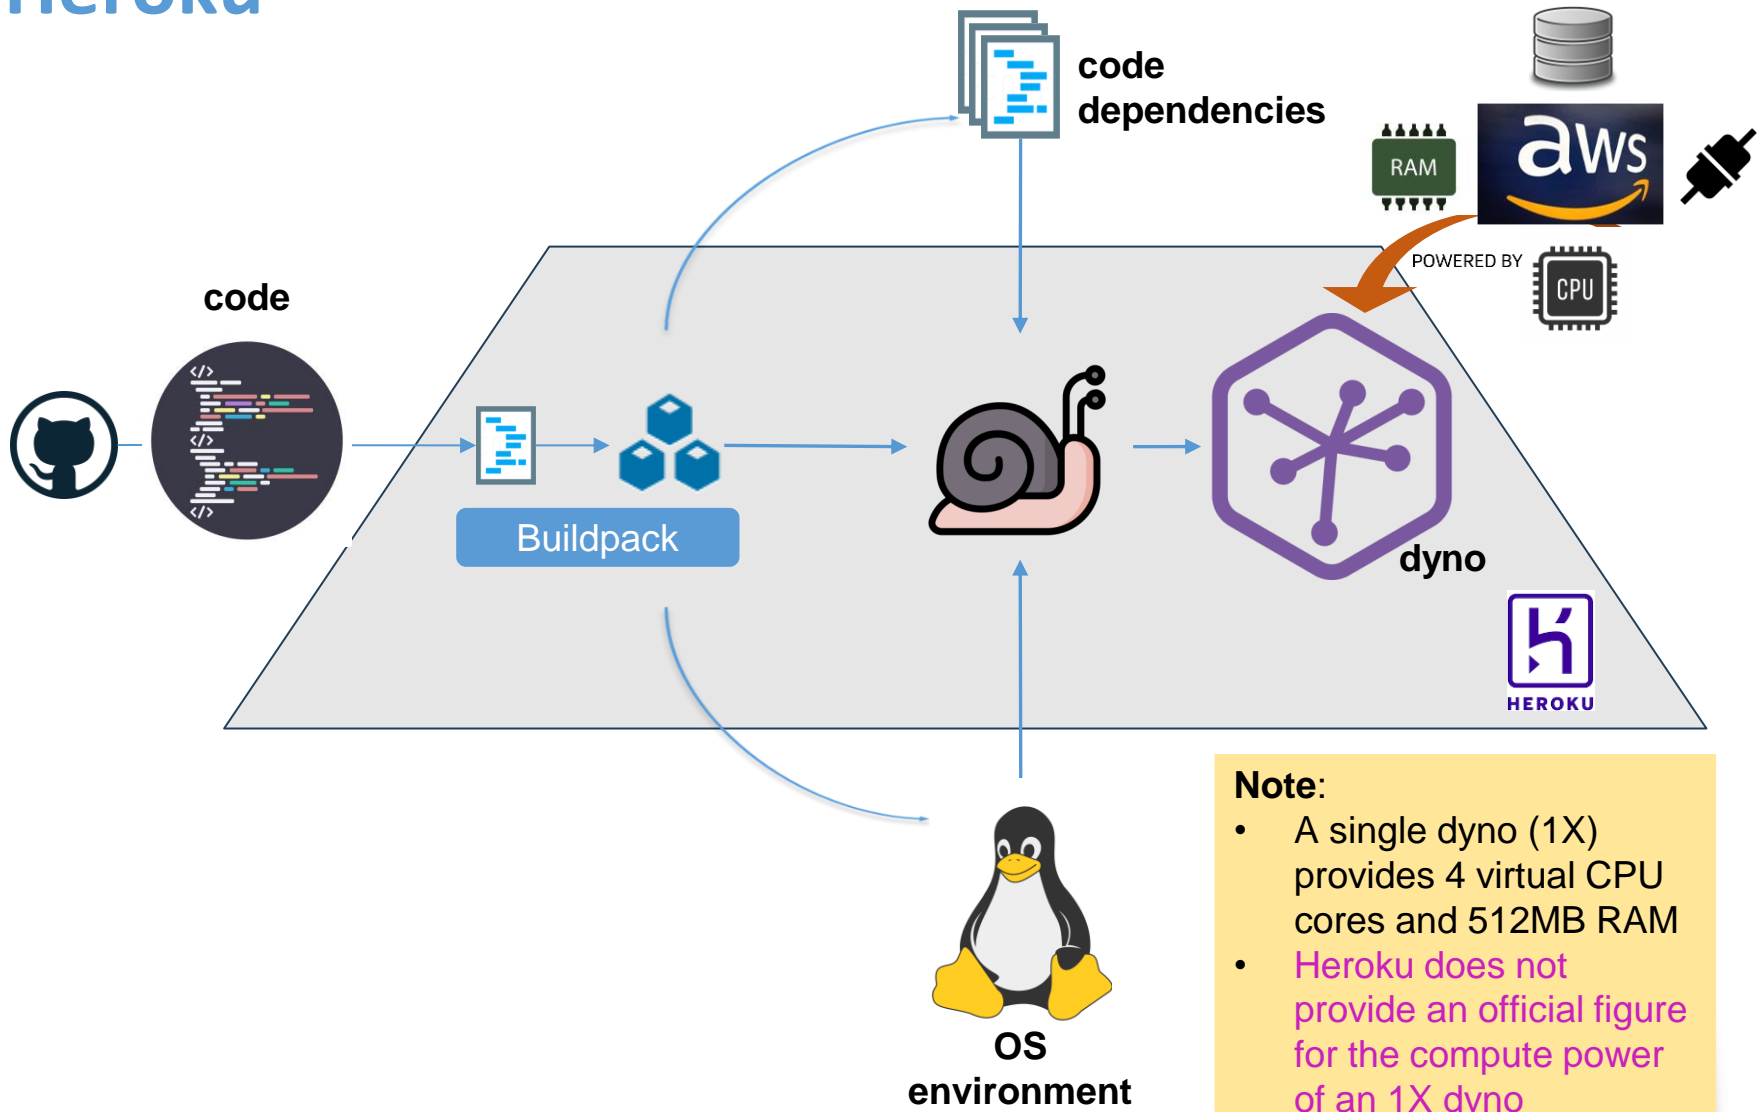
Scala



Clojure

- Provides an **operational environment** in which applications developed in the supported languages can execute
- It also provides a **development environment** to assist developers in writing applications (see later)
- But how does it work?


Heroku



Note:

- A single dyno (1X) provides 4 virtual CPU cores and 512MB RAM
- Heroku does not provide an official figure for the compute power of an 1X dyno

Heroku: Buildpacks



Buildpack

■ Core concept on Heroku platform

Provides a simple and flexible way to package and deploy applications

- Builds the application
- Sets up appropriate **runtime environment** for the application to execute

Two kinds:


Official

- Heroku provides a collection of officially supported buildpacks for popular languages and frameworks such as Node.js, Ruby, Python, Java, etc.
- Maintained by Heroku
- Designed to be used out-of-the-box for most common use cases

3rd party

- Developers provide custom buildpacks
- Customize the build and deployment process of an application

Heroku: Buildpacks



Buildpack

By invoking an appropriate runtime (e.g., **JRE** and hence **javac**)

Compilation:

Transforming a program P written in a language L to a semantically-equivalent program P' expressed in some **object code** language L'

In Java L' is called **bytecode**: a platform independent non-executable format

Java (and .NET) uses Just-In-Time (JIT) compilation: `javac` compiles source code into bytecode which is then **interpreted**

As opposed to (ahead of time, AOT) compilation, **interpretation** is performed on an instruction-by-instruction basis

AOT compilation leads to shorter execution times and JIT compilation to better portability and shorter building times

■ Building an app

Linking

- Combining multiple object files and libraries into a single executable
- Resolves references to functions and variables defined in other parts of the code
- Linking is a dynamic process merging bytecode files (.class, .jar) **on demand** (at runtime)
- Merged files are either internal to the codebase or from external repos (declared in the POM file)

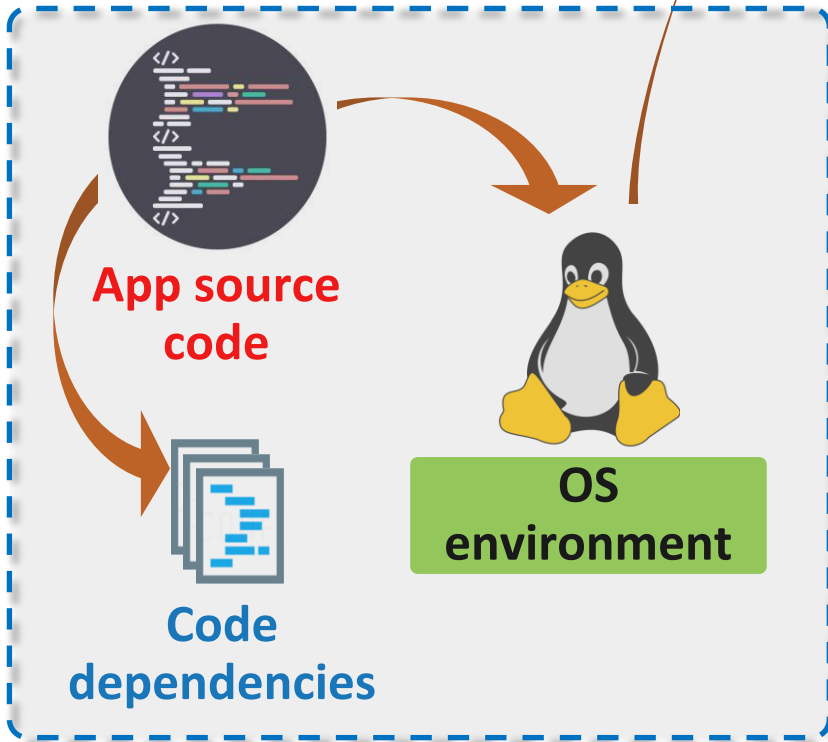
Fetches by:



Heroku: Slugs



- Bundle comprising



A **containerised** environment where the app runs

Includes:

Language runtime (e.g., JRE)

- Compiles and/or interprets source code
- Executes source code

Lightweight Linux-based OS

- Provides necessary OS runtime to support execution

Heroku: Slugs



In-slug OS provides **drivers** for accessing disks and other devices

■ **Interrupt handling**

- In-slug OS maintains **page tables** to map virtual memory space into physical
- In-slug OS manages only memory allocated to containing dyno
- Ensures **process isolation**

Of utmost importance for isolating different dynos

- In-slug OS manages **network access**
- Implements TCP/IP stack
- Opens sockets (endpoints), establishes connections, manages data transfers

Through appropriate APIs that are called from within the app

- In-slug OS implements the **upper layers** of the TCP/IP stack (http, https, app specific communication protocols)
- Lower stack layers are handled by the Heroku OS with the support of the **Heroku router**

Heroku: Router

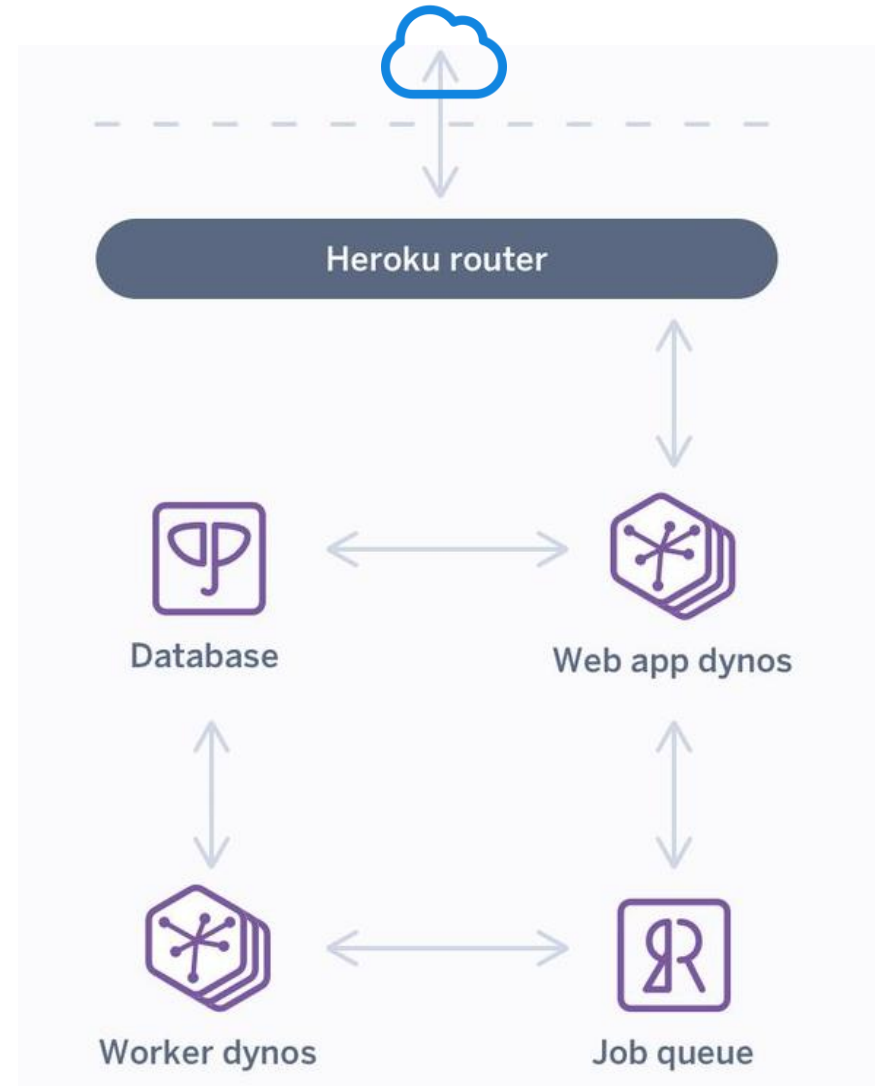
Heroku Router

Ingress traffic

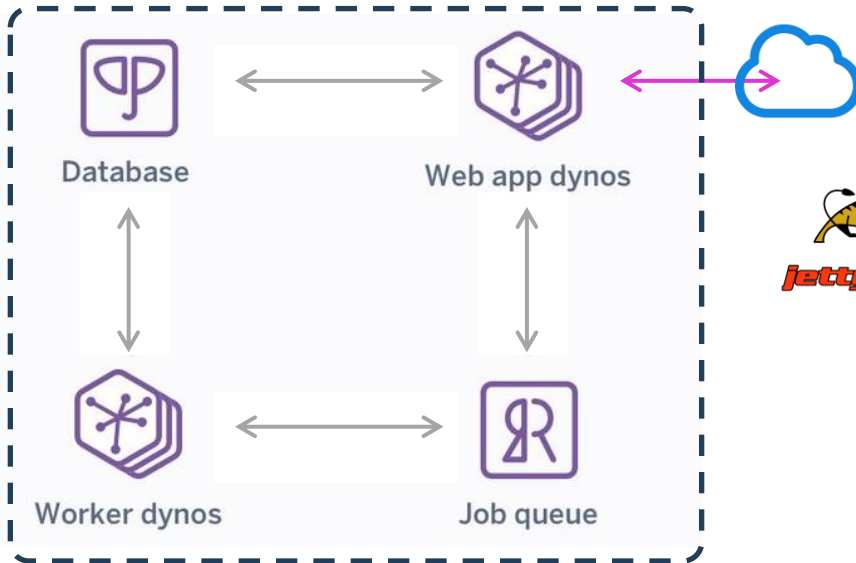
- Dynamic request routing to appropriate dyno
- Load balancing
- SSL/TLS termination (decrypting incoming traffic)
- DDoS protection
- Firewall

Egress traffic

- Routing based on IP tables
- Firewall



Heroku: Ingress and Egress Traffic



- Requests passed over to the upper layers of the TCP/IP stack




- Web server unmarshals the request which is now ready for processing




Egress traffic

- Web server marshals response
- Upper layers of the TCP/IP stack pass over response to Heroku Router

Ingress traffic




Receives incoming requests and determines the appropriate destination for each request based on routing rules and possibly **load balancing**



Receives a response and routes it towards its destination using IP tables

Heroku: Dynos



An app must always comprise at least one **web dyno**

Two kinds of dynos

Web dynos

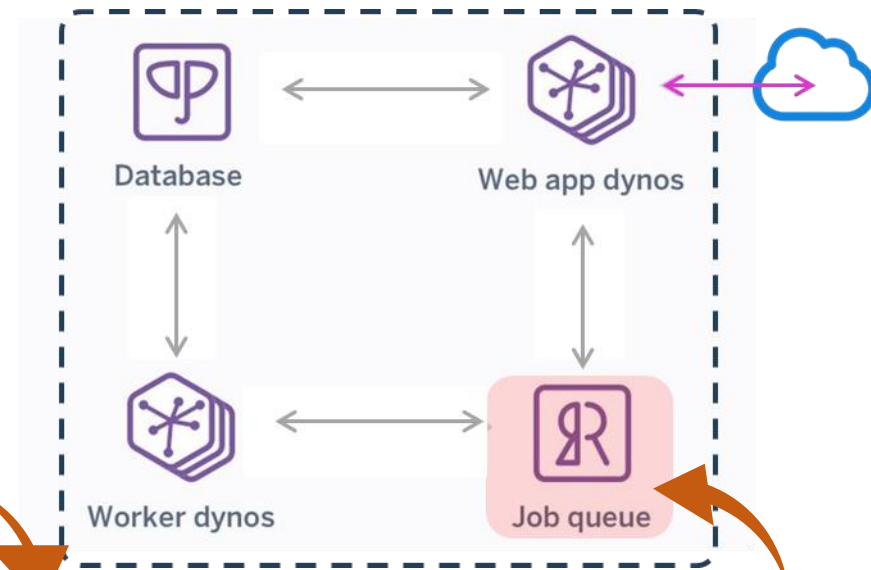
intended for serving web requests

Worker dynos

Execute background tasks

A web dyno is one whose slug includes a prepackaged webserver

Prebuilt functionality to **unmarshal** incoming requests and **marshal** responses



- No prepackaged web server
- Used for tasks that can be executed **asynchronously** and independently of incoming web requests
- Typically connect to a database

Heroku: Dynos - Procfile

A text file with commands for starting dynos

Procfiles are not necessary for simple apps (e.g., apps comprising a single web dyno)

```
web: java -Dserver.port=$PORT -Xmx512m -Dlogging.config=logback.xml -jar
target/myapp.jar
worker: java -jar target/myapp.jar -workers
scheduler: java -jar target/myapp.jar --scheduler
```

Defines three processes (dynos)

- **worker:** specifies the command to invoke a worker
- **scheduler:** specifies the command to invoke a scheduler (also a worker but this time with the label “scheduler”)

web: specifies the command to invoke a web process, on the web server at port \$PORT, the amount of RAM assigned to JVM for this app, and a custom logback

By default Heroku runs 1 web dyno automatically for a newly deployed app

Heroku: Add-ons



A modular piece of software extending the functionality of an app













Example

- Suppose a Java app that uses a database
- A popular option is to get the database as service from Heroku
- In such a case, the database must use an addon to connect to the database

Example

- Suppose an app is deployed as a web and a worker dyno
- The two dynos interact via a message queue
- The message queue may be provided as an addon

- Heroku offers a comprehensive list of add-ons
- These are either proprietary (owned by Heroku) or 3rd party

 Stackhero for InfluxDB InfluxDB on dedicated instances, up-to-date versions and attractive prices.	 Bucketeer Use Amazon S3 from your Heroku application.	 BETA RethinkDB Cloud Cloud-hosted elastic RethinkDB	 Stackhero for MariaDB MariaDB on dedicated instances, up-to-date versions and super attractive prices.
 BETA SentinelDB A privacy by design, GDPR-compliant database with per-record encryption	 Redis Enterprise Cloud From the creators of Redis. Enterprise-Class Redis for Developers (w/ Free plan)	 Crunchy Bridge Crunchy Bridge is a fully managed Postgres service from the Postgres Experts.	 Treasure Data Analytics Platform on Heroku
 Cloudcube Flexible AWS S3 file storage without the hassle.	 JawsDB Maria MariaDB, the open source drop-in replacement for MySQL now available on Heroku	 ObjectRocket for Mongo... High-Performance MongoDB + Fanatical Support™ objectrocket.com/mongodb-heroku	 Memetria for Redis** Redis 7.0, Instrumented and Scaled

and many more...

▪ Benefits

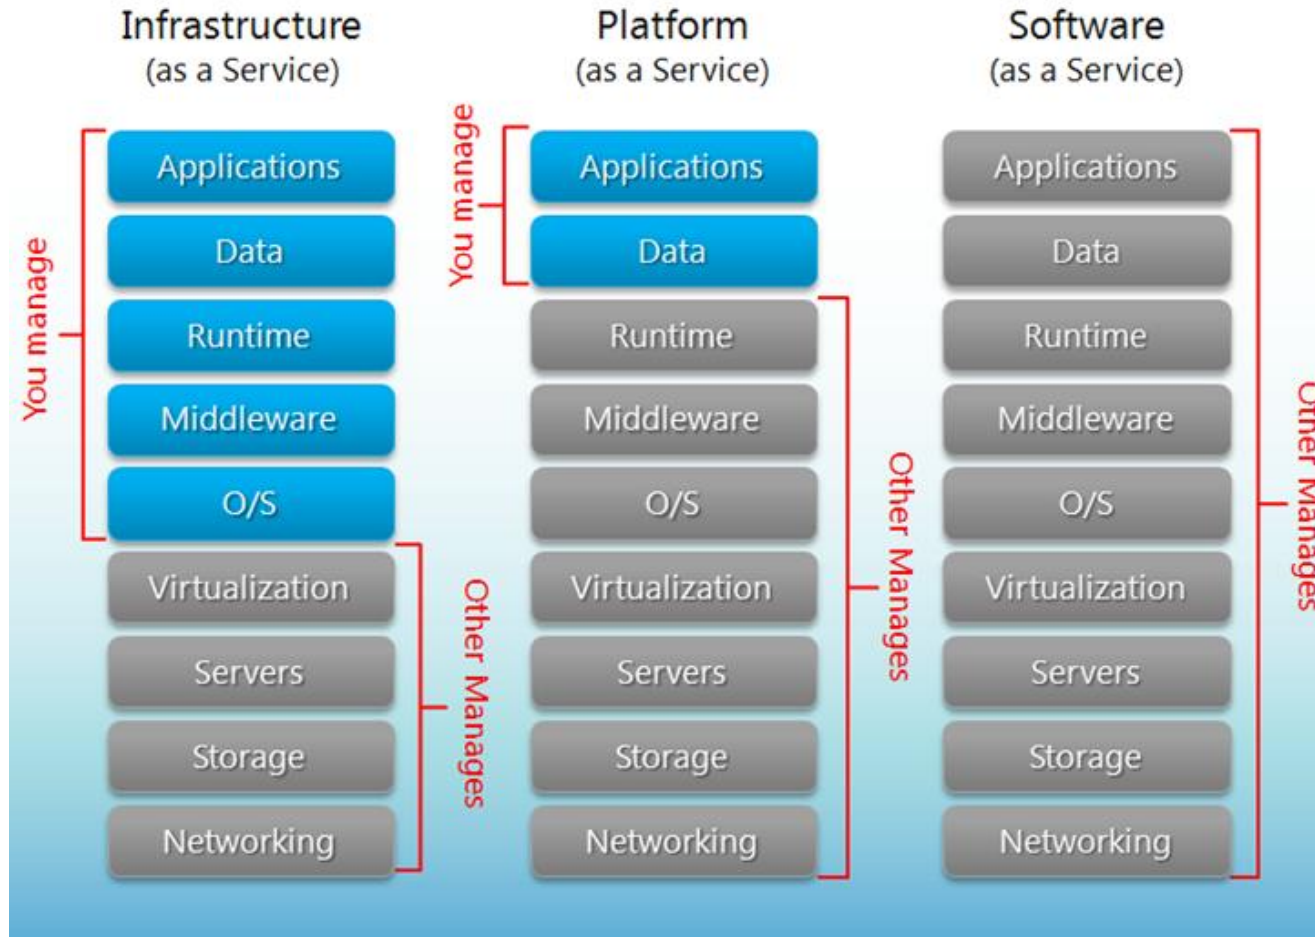
- Cost savings
 - Leveraging economies of scale allows lower infrastructure setting up, running and maintenance costs
 - The flexible pay-as-you-go model incurs lower costs as it allows usage to be monitored, measured, and billed transparently based on utilisation
- Agility and increased productivity
 - No need to spend time on planning and purchasing servers, networking equipment, etc.
 - Easy and cost-effective to quickly setup the infrastructure required for testing up new ideas
 - Development can be greatly facilitated through the use of 3rd party services
 - Faster time to market
 - Focus on innovation rather than on code deployment, infrastructure maintenance and update issues
 - Location independence

- Benefits_(contd.)
 - Performance and availability
 - Scalability to cope with increased demand loads
 - Reliability
 - Provides support for disaster recovery and business continuity
 - Power-efficient
 - Multi-tenancy and virtualisation permit less power consumption
- Criticisms
 - Security
 - Lock-in effect
 - Applications developed on a particular PaaS platform are designed to be offered on that platform
 - This also depends on the amount of 3rd party services used

- Software as a Service
- Refers to software (services and applications) that are delivered over the Internet on an on-demand basis
 - No need to go through a sales process to gain access to the software
 - The software is made available through a web browser
 - No need to install the software locally
 - Allows for the use of thin clients for accessing possibly 'fat applications' (i.e. hefty applications that wouldn't be able to run on a thin client)
 - Allows for ubiquitous access
 - The software is charged on a pay-as-you-go basis
 - The customer is charged only for the parts of the service used and for as long as they are used
 - The customer is typically billed on a recurring basis
 - The software and the entire software stack on which it operates is maintained by the vendor

SaaS

- SaaS vs PaaS vs IaaS



source: blogs.technet.microsoft.com

SaaS

- SaaS applications have a much lower barrier to entry than their locally installed competitors
- Software types that lend themselves to the SaaS model include:
 - Customer relationship management (CRM)
 - Video conferencing
 - Accounting
 - Web analytics
 - Web content management
 - Collaboration software
- SaaS applications may be used out-of-the-box
 - No need to make extensive changes or integrate with other systems
- SaaS applications are also typically offered in easily-customisable versions
 - Provide the opportunity to clients to create added value by customising the applications in the way that suits their businesses better

Note: Well known SaaS examples include Gmail, Google docs, ...

SaaS



Salesforce

SalesforceIQ Starter	Professional	Enterprise	Unlimited
Starting at \$25 user/month* (billed annually)	\$65 user/month* (billed annually)	MOST POPULAR \$125 user/month* (billed annually)	\$250 user/month* (billed annually)
Out-of-the-box CRM for up to 5 users	Complete CRM for any size team	Deeply customizable CRM for your business	Unlimited CRM power and support
TRY FOR FREE	TRY FOR FREE	TRY FOR FREE	TRY FOR FREE

* All per user products require an annual contract.

Introducing SalesforceIQ

Discover a smart, easy CRM solution built for small businesses. SalesforceIQ is an out-of-the box application powered by Relationship Intelligence that can be set up in minutes.



- Benefits
 - Cost savings
 - Leveraging economies of scale allows lower infrastructure setting up, running and maintenance costs
 - The flexible pay-as-you-go model incurs lower costs as it allows usage to be monitored, measured, and billed transparently based on utilisation
 - Agility and increased productivity
 - No need to spend time on planning and purchasing servers, networking equipment, etc.
 - Faster time to market
 - Location independence
 - High-degree of customisability
 - Performance and availability
 - Scalability to cope with increased demand loads
 - Reliability
 - Provides support for disaster recovery and business continuity

SaaS

- Benefits_(contd.)
 - Power-efficient
 - Multi-tenancy and virtualisation permit less power consumption
- Criticisms
 - Lock-in effect
 - Depends on the degree of customisation performed...